Multi-Agent Reinforcement Learning for Assessing False-Data Injection Attacks on Transportation Networks

Taha Eghtesad✸, Sirui Li❖, Yevgeniy Vorobeychik ✢, **Aron Laszka✸**

✸ Pennsylvania State University

❖ Massachusetts Institute of Technology

✢ Washington University In St. Louis

23rd International Conference On Autonomous Agents And Multi-Agent Systems (AAMAS)

Auckland, NZ, May 6-10, 2024

WHY TRANSPORTATION SERVICES ARE IMPORTANT?

PennState

- Provide access to:
  - Education
  - Healthcare
  - Emergency Services
- Contribute to:
  - Economic growth
  - Logistic services
  - Delivery of essential goods

- Disruptions can lead to:
  - Financial losses
  - Physical damage
  - Bodily harm

Images from Gemini Advanced

VULNERABILITY OF TRANSPORTATION NETWORKS

PennState

- SMS **Disinformation**
- Traffic **Sign** Manipulation
- Traffic **Signal** Manipulation
- **False Data Injection** in Navigation Applications

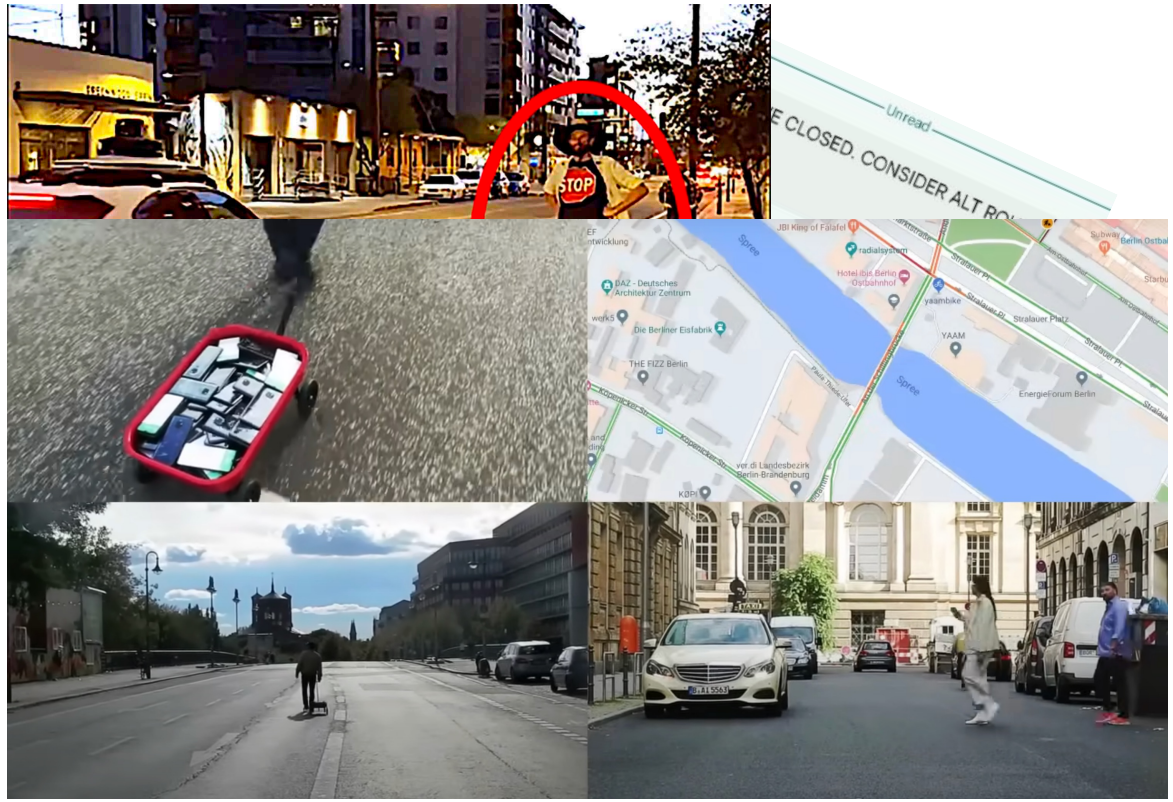

Images from https://notthebee.com/, https://9to5google.com/, and https://statescoop.com/

TRANSPORTATION NETWORK MODEL

PennState

- A **directed graph** $G = \langle V, E \rangle$ defines the transportation network's roads and intersections
- **Congestion Model**
  - Each road has a given free-flow travel time
  - The more vehicles on a given road, the higher the actual travel time
- **At each intersection**, drivers take the **shortest path to their destination** based on a navigation application

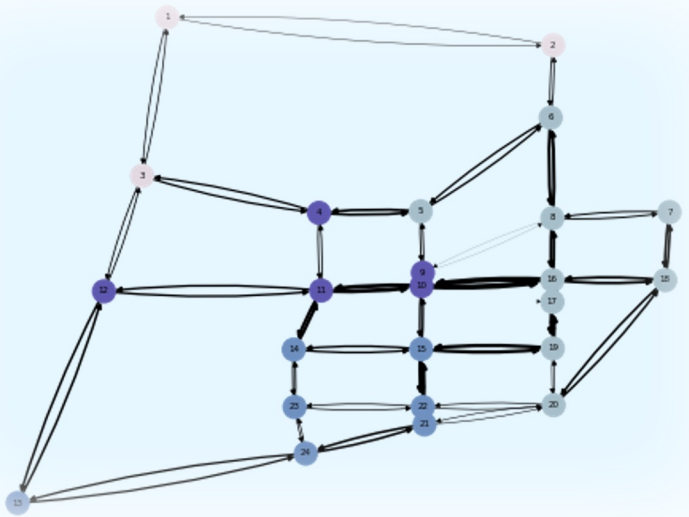

Sioux Falls, SD

FALSE DATA INJECTION (THREAT) MODEL

PennState

- The attacker has a budget to perturb perceived travel times
- The attacker perturbs perceived travel times at each step
- The drivers take a longer path due to perceived congestion

- **Strong threat model**:
  The attacker has full observation of the network
  - Vehicle locations
  - Vehicle destinations

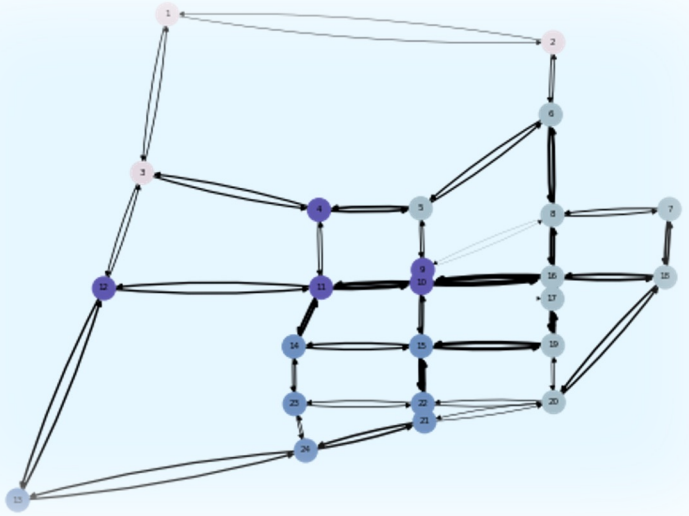

Sioux Falls, SD

PROBLEM FORMULATION

PennState

- **Assessing the extent of the damage** is the prerequisite for defense
  - An attack oracle can be used to **generate worst-case** attacks for detection and mitigation schemes
- False data injection attacks may happen **over a time horizon**
- **Uncertainty** of the environment

- The attacker can manipulate observed congestion in a navigation application
  - Restricted to a **fixed budget**
  - Able to manipulate **any road link**
  - Aiming to cause **worst-case impact**

$$MDP = \langle S, A, R, T \rangle$$

$S \mapsto$ state space
$A \mapsto$ action space
$R(s, a) \mapsto$ rewarding rule
$T(s, a) \mapsto$ transition rule

- Find a policy, mapping from **network state** to **perturbations**, that maximize **total travel time**
- Leading to: **Markov Decision Process (MDP)** formulation

REWARD, ACTION, AND STATE SPACE

PennState

- Objective
  - Goal: maximize total travel time
  - Reward: r^t = number of vehicles in traffic
- Action Space
  - Perturb observed edge travel times restricted to a budget
  - Action space: $|a^t|_1 \leq B$ and $a_e^t \geq 0$
- State Space
  - Vehicle locations and destinations

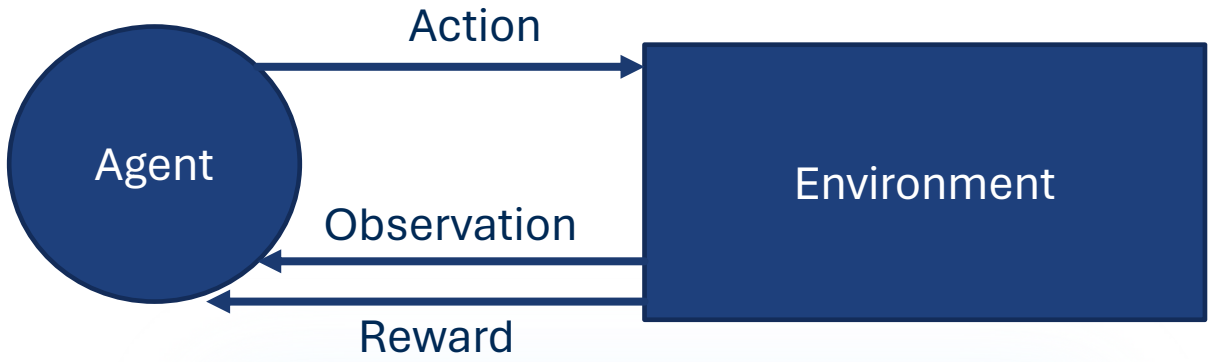

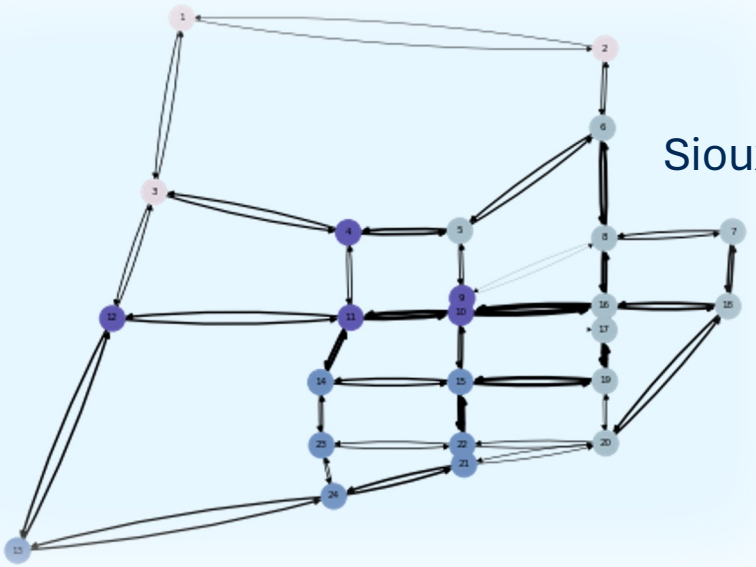

Sioux Falls, SD

DEEP REINFORCEMENT LEARNING AS ATTACK ORACLE

PennState

- Reinforcement Learning

optimize $\pi(o^t) \mapsto a^t$

max $\mathbb{E}[\Sigma_{\tau=0}^{\infty} \gamma^{\tau} \cdot r^{t+\tau} | \pi]$

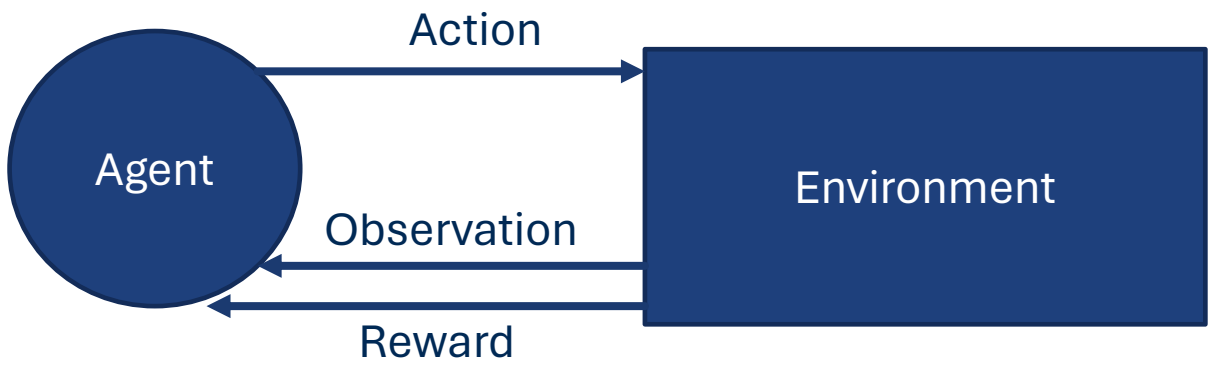

- **Critic**: $Q(o^t, a^t) \leftarrow r^t + \max_{a'} Q(o^{t+1}, a')$
  - Updated by gradient descent, reducing Mean Squared Bellman Error
- **Actor**: $\pi(o^t) \leftarrow \text{argmax}_{a'} Q(o^{t+1}, a')$
  - Updated with gradient ascent, increasing Q

FEATURE EXTRACTION

- Features for edge e
  1. Number of vehicles that are at an intersection with an unperturbed shortest path to the destination that passes through e
  2. Number of vehicles that are on an edge but will take e as the shortest path
  3. Number of vehicles that are at an intersection that will immediately take e as their shortest path without perturbation
  4. Number of vehicles currently on e
  5. Sum of remaining travel times of vehicles currently on edge e
- State represented as $|E|\times 5$ vector

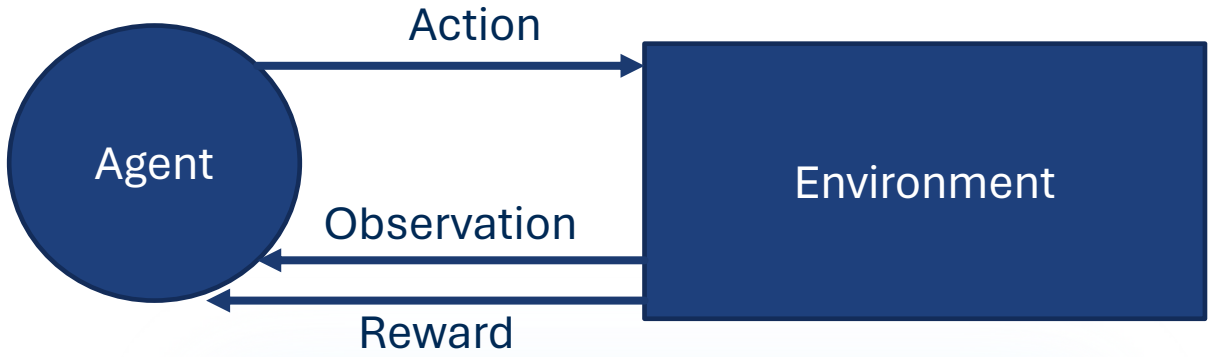

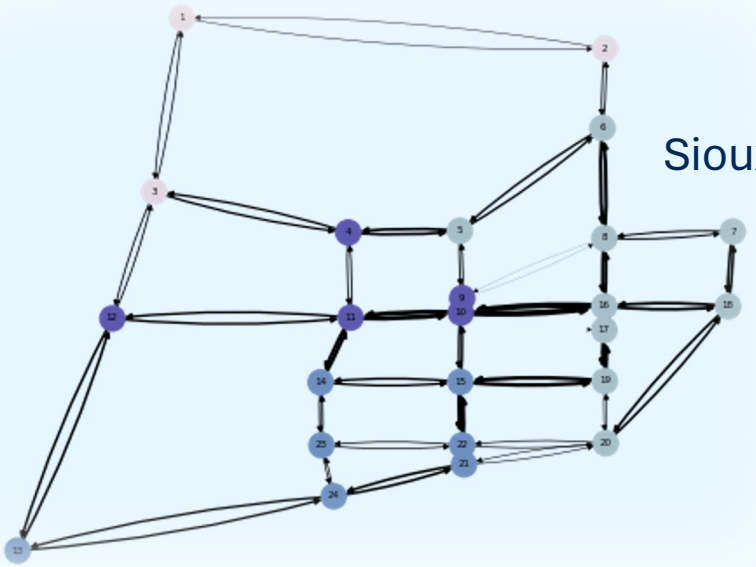

Sioux Falls, SD

ASSESSMENT THROUGH DEEP REINFORCEMENT LEARNING

PennState

- The attacker could output perturbations for **hundreds of city roads**
- General-purpose reinforcement learning algorithms (e.g., DDPG) are **infeasible** even for a small city
  - 24 nodes and 76 edges in Sioux Falls
  - Enormous action/observation space
- It requires **millions of samples** collected from the environment
- We need a robust and feasible attack oracle

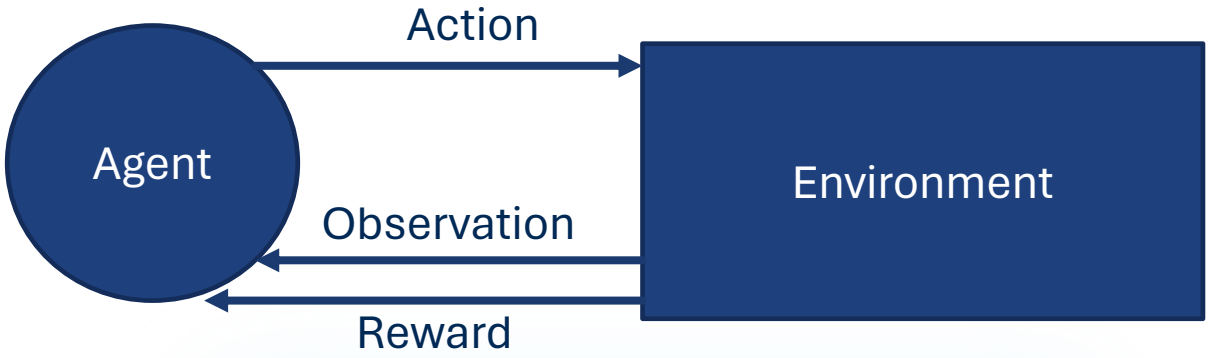

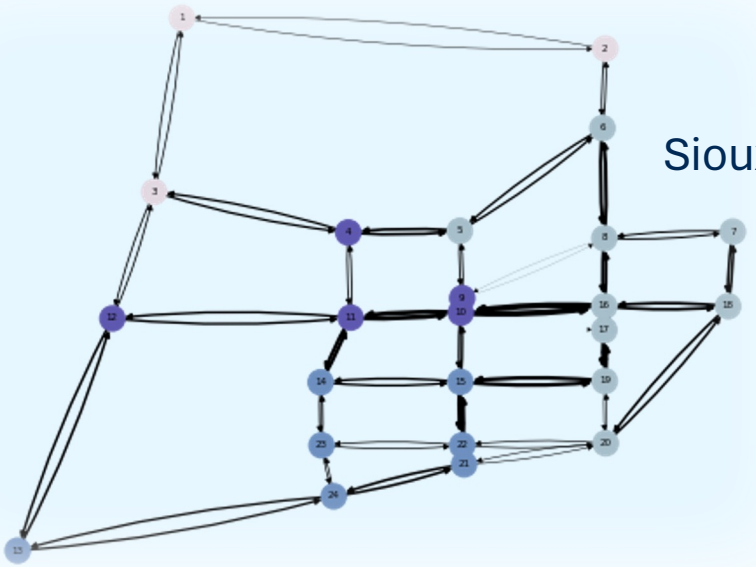

Sioux Falls, SD

HIERARCHICAL MULTI-AGENT REINFORCEMENT LEARNING

PennState

- The idea:
  - We can **divide** the network into smaller components
  - **Low-leve**l RL agents are assigned to each component
  - A **high-level** RL agent coordinates the low-level agents

- Why a high-level coordinator?
  - The total perturbations are restricted by a budget
  - Low-level agents compete over the budget

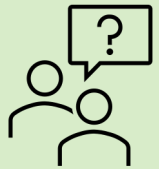

- The high-level agent allocates the perturbation budget to the component agents
- The low-level agents distribute allocated perturbation budgets to road links

HIERARCHICAL APPROACH

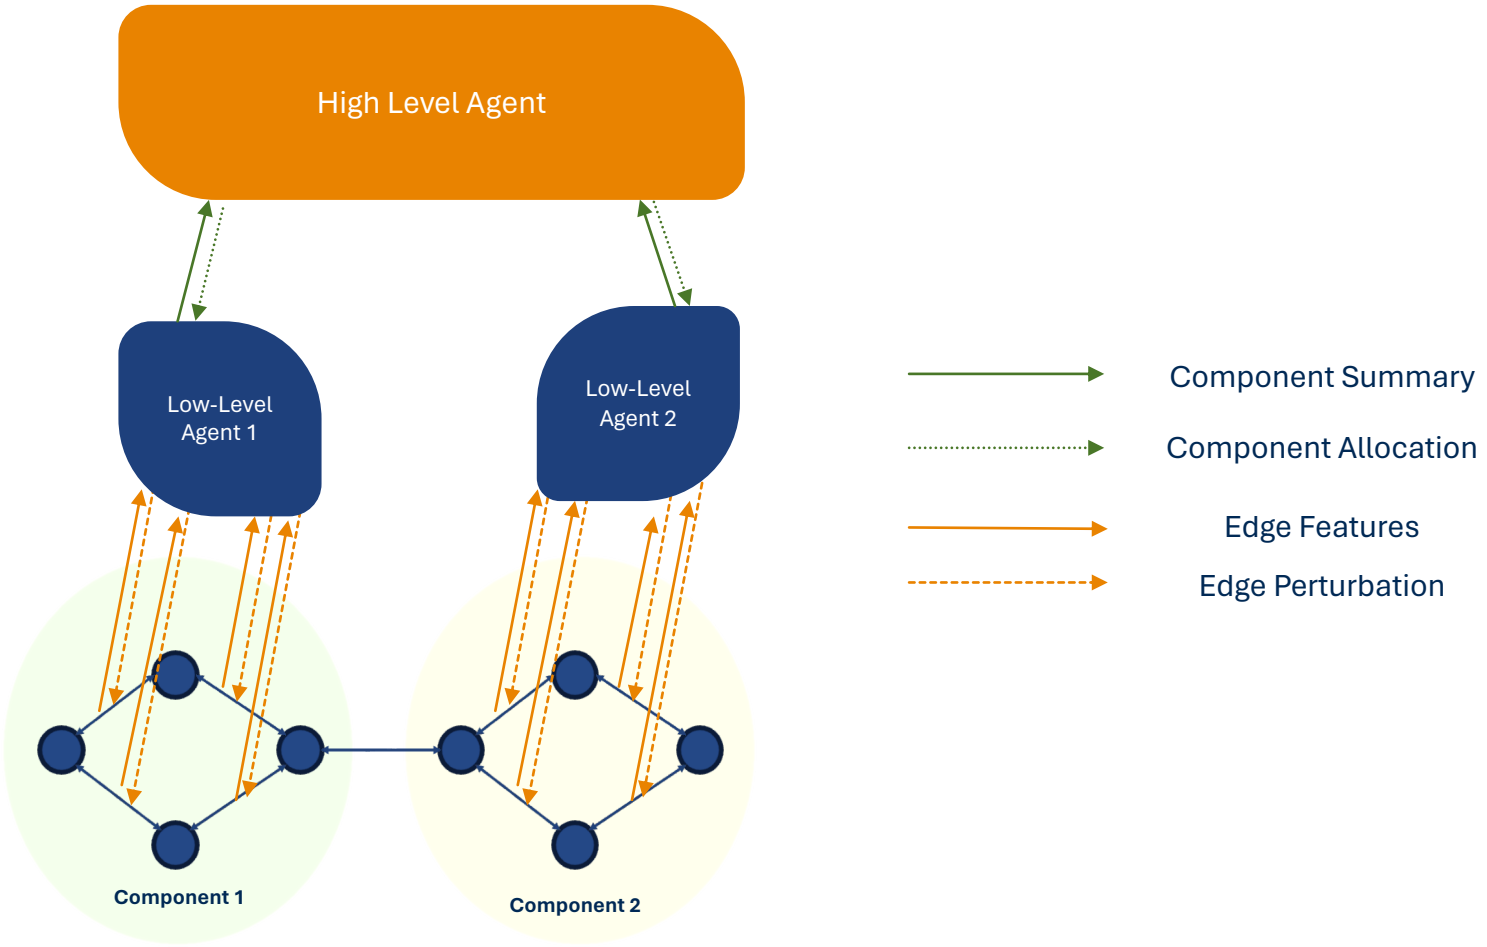

NETWORK DECOMPOSITION

PennState

Decompose the network based on **K-means clustering** by edge distance (without congestion)

Sioux Falls, SD

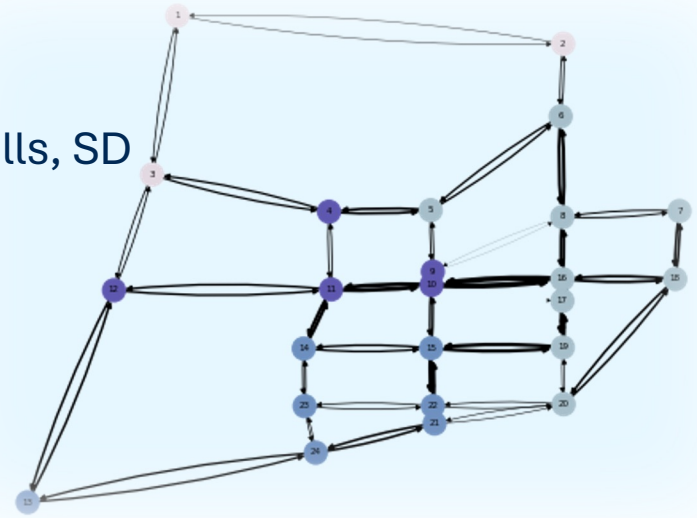

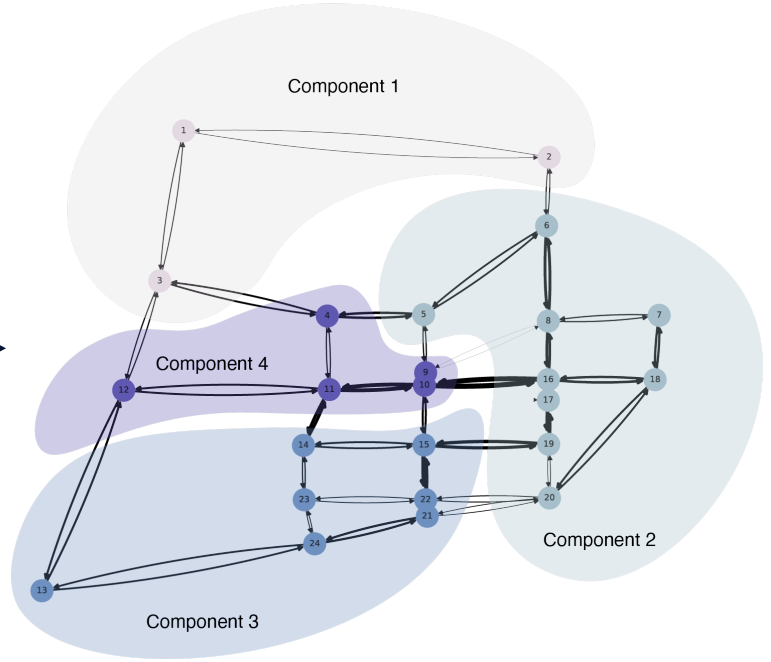

DISTRIBUTED LEARNING

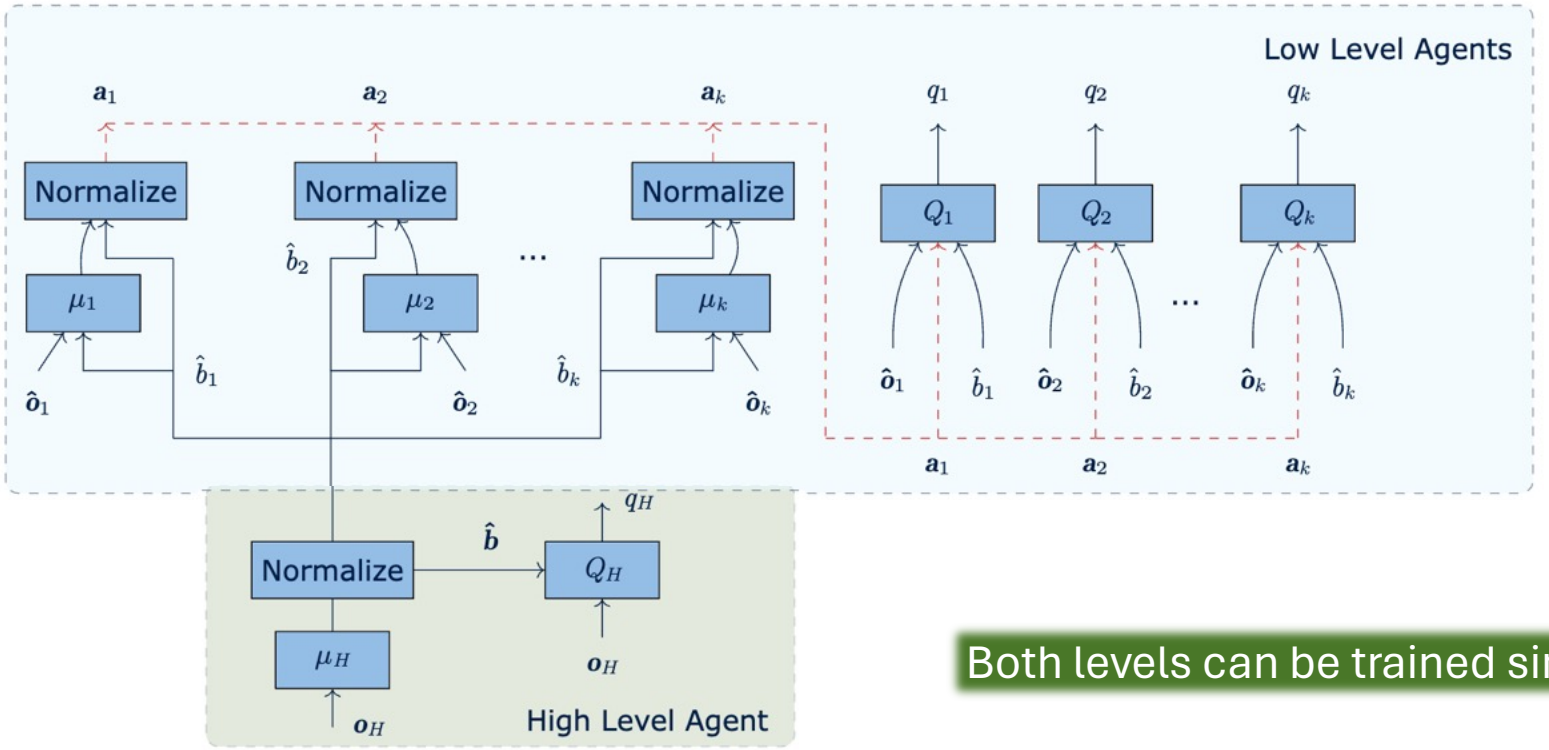

Both levels can be trained simultaneously.

EXPERIMENTAL SETUP

PennState

- Baselines
  - **Proportional** *High Level*: Allocates budget to each component based on proportion of vehicles in the component
  - **Greedy** *Low Level*: Perturbs edges by proportion of vehicles that pass through that edge
  - Random
  - DDPG without decomposition
- Hyperparameter search
  - Grid search

EVALUATIONS

PennState

Budget B = 5

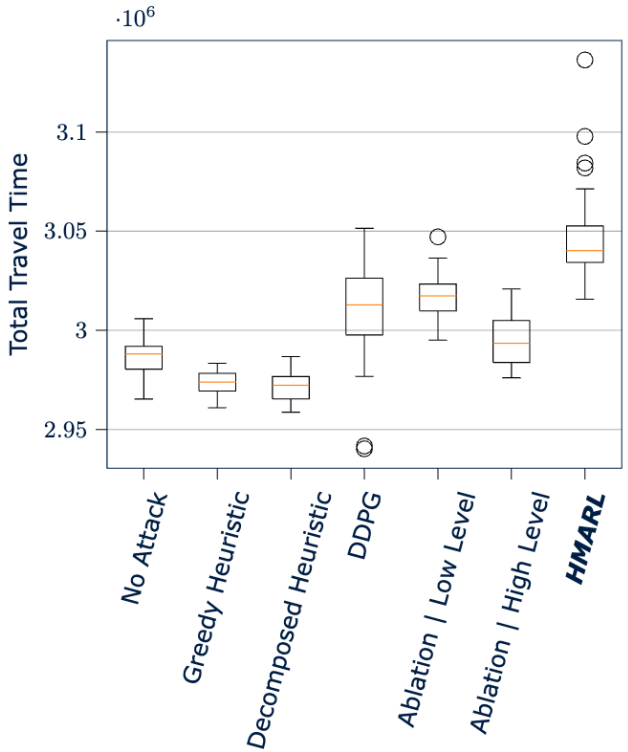

Budget B = 10

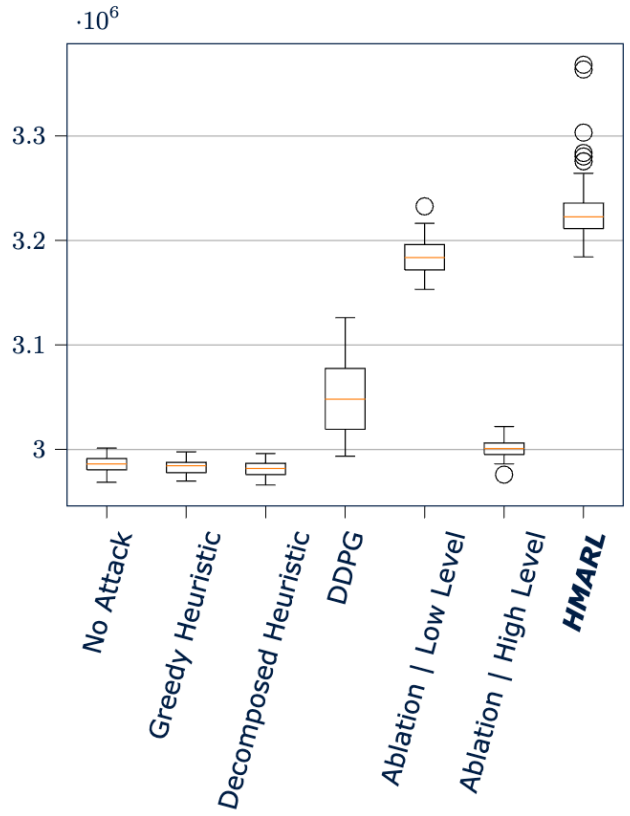

Budget B = 15

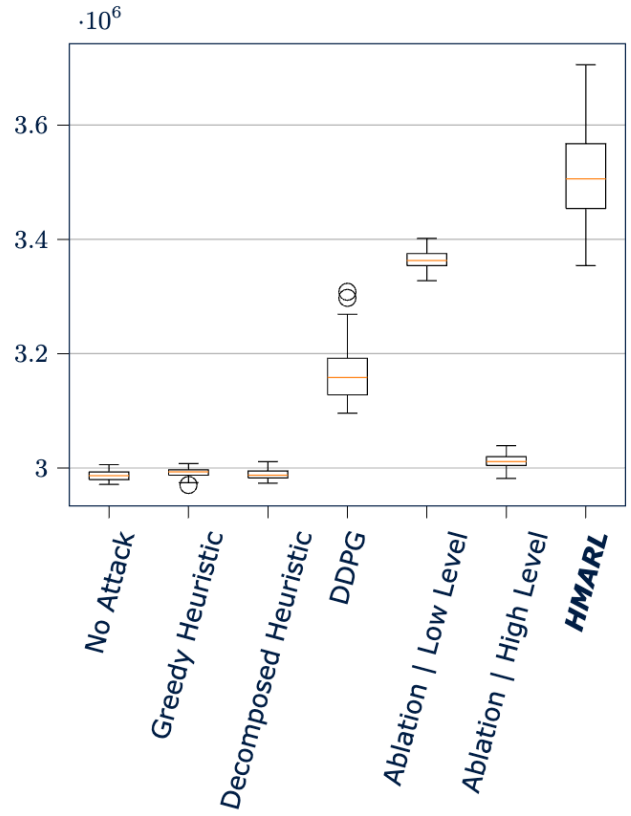

CONCLUSION

PennState

- We discussed the importance of **resiliency** of transportation networks
- We discussed how transportation networks are **vulnerable** to various attacks.
- We introduced a **model of false-data attacks against navigation** in transportation networks
- We proposed a computational method based on **multi-agent reinforcement learning** to assess against worst-case attacks
- We demonstrated the **effectiveness** of our framework on the Sioux Falls, SD benchmark network
- We showed that a **worst-case attack can increase total travel time by up to 50%**

THANKS FOR YOUR ATTENTION

Dr. Aron Laszka

alaszka@psu.edu